

PART I

Crisis

CHAPTER ONE

Baseline

Homelessness occupies a prominent place in American political life. Although less than one-fifth of 1 percent of the U.S. population experiences homelessness on a given night in the country, the issue receives considerable attention from policy makers and the general public. This spotlight is striking given the scale of the homelessness crisis when compared to other prominent social problems. That fifth-of-a-percent figure translates to about five hundred sixty-eight thousand people. To be sure, this number should feel large and unacceptable. But on an absolute basis, for example, homelessness pales in comparison to the nation's poverty crisis: Over thirty-four million Americans were living below the federal poverty line in 2019. Meanwhile, abundant evidence highlights the political preoccupation with homelessness. In 2020, a poll in Washington State revealed that voters ranked homelessness as the top priority for the state legislature—far above other common public concerns like transportation, the economy, the environment, and health care.¹ We observe a similar focus at the national level. As depicted in




Figure 1. Public interest over time for five search terms. Data source: Google Trends

Figure 1, from January 2015 to January 2020, more people in the United States searched for the term *homeless* via Google than for *inequality*, *racism*, *poverty*, and *climate change*.²

How might we explain this seemingly disproportionate interest in the issue of homelessness? Two potential explanations come immediately to mind. First, maybe this interest isn't as disproportionate as it might initially appear. That is, maybe the numbers are wrong. Among astute observers, it is well understood that official point-in-time census estimates of homelessness underestimate the true size of the population experiencing homelessness on any given night.³ For example, the federal definition excludes many precariously housed individuals and families who might be living with a friend or temporarily living in a motel room. The more expansive definition of homelessness used by the U.S. Department of Education suggests a population

of 1.35 million homeless students *without* counting their parents.⁴ Furthermore, across greater spans of time—say, lifetimes—roughly 5 percent of the population experiences homelessness at least once.⁵ In light of these figures, it is more accurate to consider homelessness as a problem that affects millions, rather than hundreds of thousands. But even the larger figure highlights the fact that only a small fraction of people living in poverty actually lose their housing.

More fundamentally, though, a second explanation for the intense interest in the topic may stem from the simple incongruity of a half million people living in shelters and on the street in the wealthiest country in the world. Reactions to this apparent paradox are diverse. For some, homelessness is a moral and political outrage indicting the capitalist system on which U.S. society is based; for others, homelessness is a scourge ruining the nation's largest and most dynamic cities. Other observers reside somewhere in the middle of this spectrum. What is uncontroversial is that homelessness elicits strong and emotional responses from all corners of society. From the perspective of the public, the intense focus on homelessness requires—and demands—an explanation. There is a strong desire to understand the causes of homelessness and where to assign blame. This book is in part concerned with the question of blame.

In January 2020, just weeks before the outbreak of the coronavirus pandemic in the United States, a long-simmering debate about the origins of and responsibility for the homelessness crisis erupted in public. Members of the federal government, including President Trump, argued vocally that the high rates of homelessness in many U.S. cities were a function of the local failings of Democratic leadership and policies. Referencing Democratic House Speaker Nancy Pelosi, the president said,

“She ought to go home and take care of her District, where the homeless is all over the place, and the tents and the filth and the garbage is eroding right into the Pacific Ocean and into their beaches.”⁶ In response to this finger-pointing, state and local policy makers—most notably California’s governor, Gavin Newsom, a Democrat—argued that a lack of federal assistance had starved local communities of sorely needed resources, and housing instability and homelessness had flourished in turn.

Certainly, some of this political jostling is a product of the polarized nature of U.S. politics in the 2020s. From voting rights to climate change—issues that would appear at face value to be resoundingly nonpartisan but which often provoke party-line votes—policy responses to (and public perception of) the issues of our time are characterized by tribalism. Tailored media narratives and the so-called filter bubbles of social media add fuel to the flame of confirmation bias. It’s harder than it should be to find fact-checked information, and it’s even harder to internalize narratives that run counter to our beliefs. In this respect, homelessness is no different: It tends to provoke hyper-partisan diagnoses and prescriptions. And as with most cases of hyper-partisanship, neither argument above—Trump’s nor Newsom’s—sufficiently explains the state of homelessness in the country. If inadequate federal support alone accounts for the crisis, why does the rate of homelessness vary so substantially across cities? Presumably, all cities would be equally starved of resources if federal retrenchment were the cause. Yet while some cities have seen rates of homelessness rise over the last ten years, many others have seen rates fall. And if Democratic mayors and governors are the problem, how can we account for the many cities and states with both Democratic leadership and policies and relatively low rates of homelessness? Unsurprisingly,

the polarized plotlines are too simple, but they draw attention to essential questions about the nature and causes of the homelessness crisis.

As Ezra Klein writes in his recent book *Why We're Polarized*, one of the other phenomena driving polarization in the country is a grafting of our political identities onto national (as opposed to local) politics.⁷ National politics, by definition, require a flattening of local variation—and in our de facto two-party system, with this flattening often comes a false dichotomization of many complex issues. This complicates the effort to respond to local issues that vary by geography—homelessness among them. In the United States, one of the most pressing and vexing questions about homelessness concerns the substantial *variation* in per capita rates of homelessness in cities across the country. Seattle and San Francisco, for example, have roughly four to five times the per capita homeless population of Chicago.⁸ The stark differences between seemingly vibrant and healthy cities invite us (and many others) to ask: Why is homelessness so bad in cities like Seattle and San Francisco? Is this a failure of individuals, politicians, markets, or other structural forces? An understanding of variation might help us unlock the drivers of this crisis.

Many of us have, for good reason, struggled to identify a credible explanation for this variation. Accounts of and references to homelessness on television, online, in newspapers, and in scholarly sources offer a long list of potential causes of the issue; among them addiction, mental illness, poverty, domestic violence, eviction, high housing costs, racial discrimination, unemployment, and many others. Reports based on interviews with people experiencing homelessness highlight a wide range of potential causes, as well. A recent report from Seattle/King County for example, noted the following self-reported causes

of homelessness among respondents to the annual point-in-time homelessness census: job loss (24 percent of respondents), alcohol or drug use (16 percent), eviction (15 percent), divorce or separation (9 percent), rent increase (8 percent), argument with family or friend (7 percent), incarceration (6 percent), and family/domestic violence (6 percent).⁹ Confronted with the question of why some cities have far greater per capita rates of homelessness than others, a reasonable, logical reaction might be to assume that higher levels of homelessness stem from higher incidences of these self-reported causal factors in these cities. In this book, we examine this logic.

While perusing any list of potential causes of homelessness, one can generally break the ostensible explanations down into two overarching categories. Some causes are individual in nature, and some are structural. The bifurcation is consistent with decades of research on poverty and homelessness. On one side of the debate are those who argue that poverty and homelessness are the result of individual factors, that vulnerabilities related to housing instability are fueled by illness, mental condition, laziness, or poor decision-making, including—for these observers—excessive drug and alcohol use. And in the central downtowns of cities like Los Angeles, San Francisco, or Seattle, thousands of unsheltered people experiencing homelessness may indeed be suffering from a substance use disorder, mentally ill, and/or unemployed. Following this logic, it is the disproportionate presence of people with these vulnerabilities in certain cities that explains the substantial variation in per capita homelessness rates around the country. Whether born in or attracted to these cities, *people* comprise the homelessness crisis, and so homelessness is an individual problem. (It is not uncommon for some to argue that homelessness is exclusively an individual choice.) On

the other side of the debate are those who argue that larger, structural forces, such as market conditions, housing costs, racism, discrimination, and inequality, causally explain the prevalence of homelessness. Under the structural explanation, homelessness is a consequence of broader and deeper societal factors driving people at the margins of society out of their housing.

Perhaps there is a middle road. The individual explanation is alluring—it's individual people who lose their housing, after all. Surely there must be systematic factors at play, though; otherwise, how could we possibly account for the dramatically different rates of homelessness around the country? Even if you were entirely convinced of the individual explanation, you would have to acknowledge that some kind of systemic variation—some combination of environmental, political, economic, and demographic trends—characterizes different places. In 2019, less than 1 in 1,000 residents were unhoused in Alabama and Mississippi, while California and Oregon had over five times that rate. Why? Existing research provides a helpful roadmap to navigate the seemingly complex and, at times, contradictory evidence about the causal drivers of homelessness. Homelessness researcher Brendan O'Flaherty, for example, suggests that to generate causal explanations of homelessness, one must consider the interaction between individual characteristics and the context in which that person resides. Either explanation alone is insufficient to explain or predict individual homelessness. By extension, he argues that people who lose their housing are effectively the wrong people in the wrong place.¹⁰ This frame helps to provide a vantage point from which to consider the central question of this book: What explains the substantial regional variation in per capita homelessness rates in the United States?

To cut to the chase, the answer is on the cover of this book: *Homelessness Is a Housing Problem*. Regional variation in rates of homelessness can be explained by the costs and availability of housing. Housing market conditions explain why Seattle has four times the per capita homelessness of Cincinnati. Housing market conditions explain why high-poverty cities like Detroit and Cleveland have low rates of homelessness. Housing market conditions also explain why some growing cities, like Charlotte, North Carolina, are not characterized by the levels of homelessness that coastal boomtowns like Boston, Seattle, Portland, and San Francisco are. Variation in rates of homelessness is not driven by more of “those people” residing in one city than another. People with a variety of health and economic vulnerabilities live in every city and county in our sample; the difference is the local context in which they live. High rental costs and low vacancy rates create a challenging market for many residents in a city, and those challenges are compounded for people with low incomes and/or physical or mental health concerns.

. . .

According to estimates from the U.S. Department of Housing and Urban Development (HUD), at least 567,715 people experienced homelessness on a single night in 2019.¹¹ But this aggregate figure masks significant geographic variation in the distribution of per capita homelessness across the country. The metropolitan areas of New York, Los Angeles, Washington, D.C., San Francisco, Seattle, and Boston alone account for over 29 percent of the homeless population in the country, despite being home to only about 7 percent of the general population. Regardless of one’s view of the problem—and the political lens through which one considers homelessness—it is reasonable to wonder what

it is about these cities that produces (or, according to some, attracts) such large and disproportionate populations of people experiencing homelessness. To explore this phenomenon, we shift the unit of analysis away from the individual and turn our attention to the metropolitan area. From this perspective, we are not interested in predicting whether a given person will experience homelessness or why someone lost their housing in the past; we are interested in understanding why, for example, the crisis is so much more extreme in Boston than in Cleveland. This analytic pivot does not preclude individual explanations for homelessness; instead, it clarifies the object in which we are interested: the city-to-city variation itself.

Understanding this variance is critical to formulating an appropriate policy response. In cities with substantial unhoused populations, it is common for rival political factions to blame one another for the crisis—a microcosm of the Trump–Newsom sparring cited above—and for the issue to devolve into a political hot potato. Often caught in the middle of this dispute are municipal leaders who are tasked with “solving” the problem (with resources that many consider to be inadequate). Societal cleavages emerge in which compassionate responses to homelessness—those that stress social service provision and respect for the dignity and rights of people experiencing homelessness—are criticized by community members who advocate a tougher response to the crisis. Proponents of the latter approach argue that overly permissive local policies have incubated an underlying problem, all while individual desperation facilitates property crime, threatens public health, and abets a deterioration of a city’s overall quality of life. The severity and polarization of homelessness is evident in public polling that identifies the problem as the highest-ranked public concern. In

the 2020 State of Washington poll of eligible voters mentioned earlier, 31 percent of voters ranked homelessness as Washington's top issue: an increase of ten percentage points over 2019.¹²

Accordingly, it is worth considering the relationship between perceptions of homelessness and its reality, not least because personal experience and anecdote play formidable roles in shaping opinions and perceptions about the issue. For housed city dwellers in Seattle, San Francisco, and Los Angeles, seeing and interacting with people experiencing homelessness is a daily occurrence. Tents dot the urban landscape; large encampments move (either voluntarily or forcibly) from neighborhood to neighborhood. There is a profound chasm in human experience in these cities, between new million-dollar condos and the tents and tarps that the unhoused use to protect themselves from the elements. And while large unsheltered populations in many coastal cities raise legitimate concerns about public safety and health—for the housed and unhoused alike—these visible reminders do not accurately reflect the homelessness problem as a whole. In most cities, the majority of people experiencing homelessness are not visible to the general population, because most people without housing sleep in shelters or other supportive housing facilities. On any given night in this country, the chronically unsheltered constitute only about one-tenth of the population experiencing homelessness. Yet the visibility—the literal conspicuousness—of the chronic, unsheltered population in many cities helps to cement a belief that people experiencing homelessness are mentally ill and/or addicted to a substance, as these conditions are disproportionately represented in the unsheltered population. Accordingly, the narrative about homelessness is often dominated by a focus on drugs and mental health, which may obscure other (often structural) explanations for the crisis.

In this book, we make an important distinction when considering the causal drivers of homelessness. First, we note *precipitating events* that can lead to a bout of homelessness. For example, in the survey of people experiencing homelessness in Seattle/King County, self-reported “primary reasons” for homelessness include divorce, domestic violence, and arguments with family or friends.¹³ As they are identified in interviews with people then-without housing, we can indeed consider these events to have produced a spell of homelessness. But we can’t consider each reason a *root cause* of a given housing crisis. If divorce is a cause of homelessness, for example, why don’t far more people lose their housing after leaving a spouse? A key point to which we return in this book is that *under certain conditions*, a range of precipitating events (like divorce) can result in homelessness—but these events ought not be considered root causes of housing instability and loss. Underlying vulnerabilities matter.

Consider the following vignette about musical chairs, often deployed by homelessness researchers, to think through causality and homelessness. We use this example to highlight the difference between a precipitating event and a root cause:

Ten friends decide to play a game of musical chairs and arrange ten chairs in a circle. A leader begins the game by turning on the music, and everyone begins to walk in a circle inside the chairs. The leader removes one chair, stops the music, and the ten friends scramble to find a spot to sit—leaving one person without a chair. The loser, Mike, was on crutches after spraining his ankle. Given his condition, he was unable to move quickly to find a chair during the scramble that ensued.

In other words, when housing is scarce, vulnerabilities and barriers to housing are magnified. Limited financial resources, mental illness, addiction, or interpersonal strife, under

a specific set of circumstances, could each precipitate a bout of homelessness—just as a sprained ankle might prevent one from finding a chair in musical chairs. But the fundamental question remains: Would we say that Mike’s ankle injury *caused* him to lose the game? Under the specific conditions of the game (say, nine chairs and ten people), Mike’s impairment prevented him from finding a chair. But under different conditions—ten chairs and ten people—Mike would have easily found one. One could argue, and we will in this book, that the fundamental *cause* of Mike’s chairlessness—was a lack of chairs, not his ankle injury. The rules of the game meant that someone had to lose.

Over the course of this book, we illustrate that personal vulnerabilities may explain *who* becomes homeless within a given community under a specific set of circumstances—but that, in aggregate, these vulnerabilities do not adequately explain regional variation in homelessness. This finding suggests that broader structural explanations of homelessness—especially those that shape housing markets—may have more explanatory power than the precipitating events frequently cited in local surveys as the “primary causes” of homelessness. Policy responses ought to be tailored accordingly. This foundation guides this project as a whole. Our central argument—that the prevalence of homelessness is driven by structural forces—is not unique in its own right. Much research has identified a causal link between housing-market conditions and homelessness. But there is little evidence that these findings have altered and shaped public perceptions about the nature of the crisis; hence our desire to package a comprehensive analysis in a single volume.

Social science research frequently relies on complex statistical methods and expansive data sources to relay a credible causal

narrative about social phenomena for the subset of readers who are trained in these methods and have access to this content (usually via university journal subscriptions). For people outside of the academy, access to this information may be limited, both with respect to the specialized nature of academic social science and the expensive paywalls of academic journals. Understanding, then, is more frequently shaped by media narratives, experience, and anecdote. Cognitive dissonance is real, too. New findings from the academy may challenge deeply held ideology, and evidence doesn't always make a difference. Accordingly, in this book, we seek to present our research by means of intuitive appeals to first principles. We use geographic variation in rates of homelessness as the foundation from which to test a wide range of potential explanations for the crisis, using an accessible analytical methodology appropriate for a broad audience. In this manner, we address many of the common narratives about homelessness in a single work with relatively simple statistical methods. Basic causal reasoning allows us to dismiss several common explanations for higher rates of homelessness: If there is no fixed statistical evidence of a positive relationship between a potential cause and our outcome of interest (i.e., rates of urban homelessness), we have to conclude it does not bear on variation in this outcome in any straightforward manner. For example, if poverty rates are low in cities with high rates of homelessness, it is impossible to attribute regional variation in homelessness to differences in the relative presence of low-income households. Applied to all potential explanations in this book, this structure provides the basis from which we ultimately attribute varying rates of homelessness to the structure of housing markets—a finding corroborated by other research leveraging different data sources and methods.

HOMELESSNESS COUNTS AND TRENDS

To study variation in rates of homelessness, we shift the unit of analysis from individuals experiencing homelessness to metropolitan areas. This is a book about cities, not individual people. We seek to explain why certain geographic locations produce (or otherwise report) disproportionately high rates of homelessness. Because homelessness is largely an urban phenomenon, we focus our attention on the largest urban areas in the United States.

The first step in understanding variation in rates of homelessness is to understand the manners in which communities measure homelessness and deliver programming. In 1987, the U.S. Congress passed the McKinney-Vento Homeless Assistance Act, which created the contemporary administrative machinery behind the federal response to homelessness. A critical component of McKinney-Vento was the stipulation that federal money was to be distributed directly to jurisdictions to fund local service delivery. To facilitate the flow of funds from the federal government to local communities, HUD required states and municipalities to self-organize into units of geographic aggregation called Continuums of Care (CoC). Today, CoCs are the main administrative entities that manage homelessness programming, allocate federal funding to local service providers, and conduct the Congressionally mandated one-night census of homelessness. Virtually every locality in the country is covered by a CoC, but the construction and distribution of CoCs varies from state to state.¹⁴ Most urban areas are covered by a single CoC, while smaller cities and rural areas might bundle together in a CoC that covers a large geographic area.

Ohio, for example, is divided into nine different CoCs (see Figure 2). Eight of the CoCs cover the most populous counties




Figure 2. Ohio Continuum of Care map. The state is home to nine HUD CoCs: eight counties and the balance of state. Data source: HUD

in the state, including Cuyahoga County (Cleveland), Lucas County (Toledo), Franklin County (Columbus), and Hamilton County (Cincinnati). The largest CoC (in a geographic sense) is OH-507, which encompasses the entire balance of the state that is not covered by one of the other eight county-based CoCs. As CoCs administer their own one-night count of people experiencing homelessness—often conducted the last week of January and known as the Point-In-Time (PIT) count—the estimated unhoused population in Ohio, in the eyes of HUD, is the sum of Ohio’s nine distinct CoC counts.

Relationships between CoC boundaries and other state and local boundaries can be messy. If we are interested in measuring homelessness in Cleveland, for example, the only geographic

unit available for analysis is the Cuyahoga County CoC, which includes both the city of Cleveland and its surrounding suburbs. Therefore, for those interested exclusively in the urban homelessness, Cleveland's homelessness rates as estimated by the Cuyahoga County CoC will be imprecise. Homelessness tends to be less prevalent in the suburbs.¹⁵ Therefore, the rate of homelessness in county CoCs is, on average, lower than it is for CoCs that cover a single city. For example, the rate of homelessness in Cook County, Illinois (including the city of Chicago) in 2019 was 1.20 per 1,000 population. In the city of Chicago alone, the per capita rate was 1.96. Because major metropolitan areas correspond to a mix of city- and county-based CoCs, in this book we compare county-based CoCs to other county-based CoCs and city-based CoCs to other city-based CoCs.

To create our study sample, we began with a list of the thirty-five largest Metropolitan Statistical Areas (MSAs) in the United States. MSAs are geographic units of at least fifty thousand people that cover a major urban center plus its surrounding areas. For each MSA on the list, we identified the primary CoC in that MSA. We then excluded six MSAs because the primary CoC in question covered too large a geographic area. (For example, the CoC that includes Houston, Texas, encompasses five different counties—and won't be useful for understanding homelessness in the Houston metropolitan area alone.) Five other MSAs were excluded using similar criteria, including Riverside, California; Denver, Colorado; Orlando, Florida; Pittsburgh, Pennsylvania; and Kansas City, Missouri. After these exclusions, our sample covers twenty-nine of the thirty-five largest MSAs in the country. We ultimately include thirty CoCs, however, because of the unique case of Cook County, Illinois—which is divided into two CoCs, one for the city of Chicago, and one for the remainder of Cook County. Chicago

is the only city in our sample with this structure. Accordingly, we include the Chicago CoC in the list of city-based CoCs, but we also separately aggregate the two CoCs to get a picture of homelessness for Cook County as a whole. We include Cook County in our list of county CoCs as well, leaving a final sample of nineteen county-based CoCs and eleven city-based CoCs. In 2019, the collection of thirty CoCs in our sample accounted for roughly 45 percent of all homelessness in the United States. We compare these regions every year from 2007 to 2019.

Returning to the regional differences that motivated this book, the following graphs offer a visual explanation of the variation in per capita rates of U.S. homelessness in the country over this time period.¹⁶ Figures 3 and 4 show the per capita rates of homelessness in the city and county CoCs in 2007 and 2019—the beginning and ending years for our sample period.

As the figures illustrate, variation in rates of homelessness is not a new phenomenon: The 2007 figures show similarly wide-ranging dynamics as those from 2019. In our sample, we see single-night rates of homelessness anywhere between about 1 and 10 unhoused people per 1,000 population. The second key takeaway from these figures is that, generally speaking, high per capita locations in 2007 also saw high rates in 2019. Washington, D.C., New York City, Boston, San Francisco, King County (Seattle), Multnomah County (Portland), Los Angeles County, and Santa Clara County (San Jose) have persistently seen the highest rates of per capita homelessness over the thirteen years covered in this book. We are interested in what differentiates these cities from others. To the extent that policy choices, macroeconomic trends, or local cultural factors may drive variation in rates of homelessness, we want to know what differentiates Multnomah County, Oregon, from Maricopa County, Arizona.




Figure 3. Per capita rates of homelessness in select U.S. regions, 2007. Dashed lines indicate city and county averages of per capita PIT counts. Data source: HUD

Throughout this book, to supplement our analyses of per capita homelessness in cities and counties, on occasion, we also deploy a simple indexing approach that allows us to compare city and county CoCs directly. To create an indexed value of homelessness intensity, we divide each measurement of city per capita homelessness by the largest observed city rate across the years in our sample (2007–2019), divide each measurement of county per capita homelessness by the largest observed county rate, and then combine the transformed values into a single measure. Doing so allows us to get a sense of how the severity of a given city’s or county’s homelessness crisis evolves over time, relative to other cities and counties. The indexing approach is




Figure 4. Per capita rates of homelessness in select U.S. regions, 2019. Dashed lines indicate city and country averages of per capita PIT counts. Data source: HUD

a kind of ranking function. It's not superior to our bifurcated approach to presenting city and county rates; it complements it.

Using indexed values, Figure 5 below provides a comparison of the relative ranks of indexed rates of homelessness in each CoC in 2007 compared to 2019. (The CoCs with observations excluded in 2007 are removed from the analysis.) We observe some modest movement in the rank ordering of cities, but generally speaking, regions with high per capita homelessness in 2019 also had high rates a decade earlier.

Over the course of the book, we also make use of some core concepts from statistics to illustrate key points. The first of these is the *median*, a summary statistic that indicates the midpoint in a distribution of values. The median may differ meaningfully from




Figure 5. Rank of indexed rates of homelessness, 2010 v. 2019. Vertical position indicates rank of indexed per capita PIT counts for select U.S. regions. Data source: HUD

the average or *mean* of this distribution. Accordingly, we can understand the median as helping us understand the shape of a distribution in a manner that is less sensitive to large outliers—which drag up or down the mean. To further understand the shape of these distributions, we also measure a quantity known as the *variance*. The variance of a distribution measures how dispersed a set of values are around their mean. Mathematically, variance corresponds to the squared *standard deviation* of a distribution, a similar measure assessing dispersion. Small values for standard deviations and variance correspond to narrow distributions, while large values correspond to wide distributions. In this book, we use the word *variation* to mean “differences between measurements,” while we use *variance* to indicate the mathematical quantity just described. Generally, the book attempts to

account for variation between regions by examining simple statistical models and evaluating the degree to which they explain the variance of the distributions in question.

Consider an example in which we analyze the relationship between age and height among children aged eighteen and under. The vertical axis on our chart measures height and the horizontal axis measures age. We can quantify this relationship using a scatterplot of dots, in which each dot represents one person and illustrates their height and age. After placing all dots in our data set on the scatterplot, we can assess what kind of relationship exists between the variables and how much of the variation in height can be explained by age. Because children become taller as they age, we might expect to see an upward sloping cloud of dots, but we probably wouldn't expect all the dots to fall along a perfectly straight line. Instead, we'd observe some variation. Some people are short, some are tall, some grow early, and some grow later. We can use statistics to measure the amount of variation in one variable (height) that's captured by variation in the other (age).

To do so, in several graphics throughout the book, we deploy a statistic known as the coefficient of determination, which for unfortunate mathematical reasons goes by the abbreviation R^2 . This quantity (pronounced "R-squared") offers an estimate of the amount of variance that we might consider to be captured—that is, explained—by a line drawn through the scatterplot of points. In particular, we'll draw a line through the points that minimizes the total vertical distance between all the points on the plot and the line itself. That exercise represents a *linear regression*—a statement about one variable in terms of another, as characterized by that best-fit line. (The formula for calculating R^2 subtracts the proportion of variance *unexplained* by

this best-fit line from the number 1—leaving the proportion of *explained* variance.) R^2 tends to vary between 0 and 1, with values closer to 1 indicating a greater proportion of explained variance. There’s no hard rule governing which values of R^2 imply small or large amounts of explained variance, but generally we might say that values of R^2 below 0.1 indicate very little explanation, while values above 0.3 indicate much stronger explanatory relationships. That is, it’s important to note that R^2 doesn’t tell us everything about these relationships. For example, on its own, it won’t help us separate correlation from causality, it won’t tell us if we’re missing any important variables in our statistical model, and it won’t tell us if we have enough data to draw solid conclusions. Nonetheless, it’s a useful indicator of the coupling between two variables. Returning to our example of the relationship between age and height, it is likely that the R^2 would be high—age is a strong predictor of height among children—but it wouldn’t be 1.0. There still exists plenty of variation in height among children that cannot be explained by age.

While the story about homelessness in major metropolitan areas has been generally consistent since 2007, some critical trends have emerged—see Figure 6. First, at a national level, overall levels of homelessness have fallen over this period—and this trend is apparent in the thirty CoCs in our sample.¹⁷ While overall homelessness has fallen, the variance between different cities’ rates of homelessness has increased. In practice, that means that while falling at a national level, homelessness has become increasingly concentrated in a few cities over this time period. Given the uneven progress toward reducing levels of homelessness, the task of understanding the drivers of regional variation is cast in an important light. If people and cities experience homelessness at increasingly different rates, it’s worth asking why.




Figure 6. While median per capita homelessness has decreased over the last decade, its variance has increased. Dashed lines indicate linear regressions of year and median per capita PIT counts (and their variance) between 2007 and 2019 for a sample of U.S. regions. Data source: HUD

Over the course of writing this book, two major social events occurred with direct implications for understanding the issue of homelessness in the country. First, in March 2020, as the COVID-19 outbreak swept across the United States and residents went into lockdown, businesses and schools closed in the most immediate cessation of economic and social activity on record. The virus represented a particular concern for the most vulnerable in society, including elders, those with underlying health conditions, and people without permanent housing. Accordingly, many jurisdictions took extraordinary steps to protect the health of people experiencing homelessness—and limit community spread of the virus—by moving portions of their homeless population to hotels and motels.

In summer and autumn 2020, I (Gregg) partnered with colleagues from King County and the University of Washington to

evaluate the region's COVID-19 homelessness response. In the early days of the pandemic, the City of Seattle, King County, and their partner agencies moved over seven hundred people from congregate shelter settings into hotel rooms throughout the county—with the primary aim of adhering to public health guidelines and keeping a highly susceptible population safe. Our research—which included an analysis of quantitative data on infection rates, housing stability, exits to permanent housing, and 911 dispatch calls; as well as qualitative data from interviews with shelter residents who moved to hotels and agency staff—showed that the intervention wasn't just successful at limiting the spread of COVID-19, but also in producing other benefits on various measures of well-being, including improved health, better sleep, feelings of safety and security, less interpersonal conflict, and greater and more optimistic focus on the future (including education, employment, and permanent housing).¹⁸ These results, in part, prompted King County Executive Dow Constantine to propose an additional sales tax to fund the purchase of hotels as supplementary housing for people experiencing homelessness.¹⁹ In other words, in the Puget Sound region, out of crisis has come a major public investment in housing solutions for people experiencing homelessness.

The economic consequences of COVID-19 have also increased the risks of falling into homelessness for many precariously housed people around the country. Early estimates suggest that tens of millions of households could lose their housing as they struggle to make monthly rental payments due to loss of income.²⁰ In the early months of the crisis, many states and local jurisdictions enacted eviction moratoria to prevent people from losing their housing due to an inability to pay. In September 2020, the Centers for Disease Control and Prevention (CDC)

announced a national eviction moratorium, leveraging its broad authority to control the spread of the pandemic. As these eviction restrictions lapse, however, housing researchers and advocates offer dire predictions of a sharp rise in homelessness, since even under the CDC rules, renters are expected to pay any rent deferred under the moratorium.²¹ Beyond the obvious public health consequences of COVID-19, this crisis highlighted the natural interconnectedness of larger, structural forces and the dynamics of homelessness.

In the midst of the country's early grappling with the novel coronavirus, the May 2020 killing of George Floyd in Minneapolis at the hands of the police ignited global protests of police violence and racial inequality. In the United States, these protests brought structural and systemic racism to the forefront of public and political discourse—and created the newest opportunity for substantive movement toward dismantling structural racism in the country. Racism is central to discussions that aim to reveal the causes and consequences of homelessness. Given their representation in the general U.S. population, Black, Native, and Hispanic/Latino individuals and families are disproportionately represented in the homeless population.²² While Black people, for example, make up about 13 percent of the U.S. population, HUD reported to Congress in 2019 that almost 40 percent of the population experiencing homelessness was Black.²³ This fact alone ought to be unsurprising: A knot of conspicuous, racialized structural disadvantages—in housing, banking and lending practices, education, health care, employment, and policing and incarceration—readily amplify homelessness risk.

What's essential in these two extraordinarily salient crises—the coronavirus pandemic and the latest reckoning with structural and systemic racism—is the manner in which they

highlight whom we're talking about when we talk about homelessness risk. To the extent that this is a book about regional variation in homelessness, it is also a book about deck-stacking. Place is where it happens. If we want to understand the factors that cause people to lose their housing, we need to understand why and how those forces vary from city to city. For many, the pandemic and the Black Lives Matter movement have driven home the interconnectedness of U.S. society, not least because the health of one's neighbor is directly related to the health of oneself. But just as racism is not a monolith, homelessness is multifaceted, and households' experiences with housing vary with identity and geography. Here, we argue that an effective policy response to homelessness will only come from acknowledging and responding to these differences.

APPROACH

This book is split into three brief sections. We conclude the first section in the next chapter, in which we lay out the current state of knowledge about homelessness in the country. We provide an overview of existing academic social-science research on the topic and offer descriptive statistics about homelessness. The purpose of the second chapter is to place the reader in a position to engage critically with the specific causal arguments presented in the second section of the book.

In part 2, we consider the various potential explanations for the substantial variation in per capita homeless populations around the country. In chapter 3, we analyze a range of individual and household vulnerabilities and attributes and ask whether these common narratives explain regional variation. These data and analyses provide compelling evidence that the answer to that

question, most simply, is no. The homelessness crisis in coastal cities cannot be explained by disproportionate levels of drug use, mental illness, or poverty. In chapter 4, we then consider local culture and context, analyzing how variations in weather, local political climate, the mobility of low-income households, and the generosity of local welfare provision may influence rates of homelessness. Similarly, we find that these common explanations do not account for observed regional variation. Finally, we consider a third category of potential explanations: housing market conditions. In chapter 5, we consider housing costs, housing cost burdens, and housing availability as candidate explanations for intercity variation. In this analysis, two explanations emerge as credible factors: absolute rent levels and rental market vacancy rates. We argue that after eliminating a wide range of potential explanations for the variation in question, the descriptive and correlative findings in chapter 5 together offer the most compelling explanation of regional variation in rates of homelessness.

In part 3 of the book, we synthesize our findings in two policy-oriented chapters. In chapter 6, we propose a typology of cities to explain why certain cities—certain types of cities—experience elevated rates of homelessness; while other cities, relatively speaking, do not. Combining our data with principles from the field of urban economics, we construct a framework that, importantly, helps us understand why high-growth boomtowns don't always see significant rates of homelessness. Charlotte, for example, has grown as fast as San Francisco and Seattle, but because of a relatively robust housing supply response, the city has not faced the housing shortages that plague many coastal cities. The typology also demonstrates how population declines help to explain why a large, vibrant city like Chicago has relatively low rates of homelessness: A falling population

in Chicago has created higher rental market vacancy rates and lower prices, which produces a more accommodating housing market (relatively speaking) for vulnerable households.

In chapter 7, we conclude by presenting a broad proposal to end homelessness in the United States. Ultimately, in the long run, the prescription is simple: Policymakers must increase the number of affordable housing units and provide subsidies and rental assistance to households to ensure they can access housing. In the short run, competing demands and a lack of resources makes decision-making more challenging. Local jurisdictions must balance the needs for a more robust crisis response (i.e., greater emergency shelter capacity) with the desire to increase the supply of affordable housing. In reality, cities must devote resources to both of these responses.

To create a sustainable, robust response to homelessness, we argue that three interrelated steps are required. First, public perception of homelessness must change. As long as we continue to frame homelessness as an individual problem, we will struggle to make the structural investments needed to end it. Second, this crisis requires far greater resources from all levels of government. Existing investments, while substantial, are insufficient given the scale of the problem. Last, we encourage a broader systems approach to addressing homelessness. Focusing on three stages of the system—inflow, crisis response, and outflow—are necessary to move people out of homelessness and into stable, permanent housing. A lack of focus on any one of these stages will produce a system out of balance—and high levels of homelessness will persist.

Finally, a word on the motivation for this book. Both Gregg and Clayton are engaged in the study of and response to homelessness in the Puget Sound region. As one of the areas of the

country most affected by this crisis, understanding what drives homelessness in our region is a topic of great civic importance. In the years leading up to writing this book, we have been amazed that—despite our community wrestling with homelessness for many years—there is a lack of general understanding about the nature and causes of homelessness. Numerous narratives compete for the public’s attention and, as a result, there is no consensus about the root causes of this crisis. Without a common understanding, it is impossible for elected leaders and the community at large to marshal the resources needed to end homelessness in our community. Much of the money spent on homelessness today constitutes a *response* to the crisis rather than an *alternative* to it. In the concluding chapter we share our vision—informed by thought leaders from around the country—for community-wide approaches that are required to prevent and limit homelessness. According to the United States Interagency Council on Homelessness, “An end to homelessness means that every community will have a comprehensive response in place that ensures homelessness is prevented wherever possible, or if it can’t be prevented, it is a rare, brief, and one-time experience.”²⁴ Fair enough. But without wrapping our head around the root of the crisis—its beginning—it’ll be difficult to find its end.